

RAI Risk Taxonomy 2.0

Audited Master Release Technical Report

Human-review recovery and independent audit release

29 August 2026

Abstract

This report documents two stages of work. First, the second-round human review instructions (an 808-row register with 251 verbatim reviewer comments) were re-applied row by row to the 798 baseline L4 cards, producing a 777-card recovery release. Second, an independent audit pipeline (AC-01 through AC-08) was applied to that release. Every judgemental step was performed by two independent reviewer agents with consensus-only application, covering restoration of hard-rule violations, absorption of semantic duplicates, MECE curation of every L3 category (duplication and attribution strength), decomposition of composite cards, and literature-grounded reinforcement of under-populated categories. No EM or Hybrid EM was re-run at any stage. **The final release contains 629 L4 cards: General 487, Agentic 65, Physical 77. All 46 L3 categories hold at least three cards and no Others assignments remain.**

1 Recovery principles

Valid edits and user-approved merges from the baseline commit were preserved. Only the deletion, merge, split, reassignment, and definition-revision instructions of the second-round review were applied, traced through `source_row_id` lineage. No L3 was created or modified; the L3 master SHA-256 hash is fixed. A separate fidelity audit of all 251 reviewer comments judged 246 applied, 3 partially applied, and 2 not applied; every affected card carries the verbatim comment and its application result.

2 Applied results

3 Audit corrections (AC-01 to AC-08)

Following a line-by-line audit of the recovery release and user approval, the corrections below were applied. All judgemental decisions are recorded in the paired reviewer ledgers; only consensus (intersection) verdicts were executed. All added or revised cards carry verified literature references (XSTest, OR-Bench, AgentHarm, Cao et al. CCS 2019, NIST SP 800-82r3, NIST SP 800-193, Sclar et al. ICLR 2024, Huang et al. ICLR 2024, Moore et al. FAccT 2025). The self-harm/suicide family remains separated into four axes: preparation, execution, encouragement, and crisis-response failure.

Action	Review rows
Reassignment	80
Rewrite and keep	39
Split	19
Merge	14
Delete	14

Table 1: 166 approved human-review actions (recovery stage)




Figure 1: Final L4 cards by domain

4 Release composition

The final release ships one L1 CSV, one L1/L2/L3 CSV, and three domain L4 CSVs. Public CSVs carry 28 columns: bilingual hierarchy definitions, facet, act-type, the verbatim human-review comment and its application result, and References; full-column originals including scoring and keyword machinery are preserved in the handover package and archive. The validation folder retains the 808-row instruction register, lineage edges, deletion tombstones, merge/decomposition/c logs, and the audit correction log.

Correction	Content	Cards
AC-01	Restored suicide preparation/execution separation	777→778
AC-02	Absorbed 58 consensus duplicates across 8 large L3s	778→720
AC-03	Annotated verbatim reviewer comments and application results	720
AC-04	Stripped scoring/keyword columns from public CSVs (full originals archived)	720
AC-05	Populated the two empty L3s with referenced cards	720→725
AC-06	Reinforced two under-populated L3s; absorbed 16 governance/manipulation duplicates	725→711
AC-07	MECE curation of all L3s: 64 absorbed, 18 reattributed, 2 removed	711→645
AC-08	Decomposed 18 consensus composite cards	645→629

Table 2: Audit correction summary




Figure 2: Card count across audit corrections