

# RAI Risk Taxonomy 2.0

## Audited Master Release Technical Report

Human-review recovery and independent audit release

29 August 2026

### Abstract

This report documents two stages of work. First, the second-round human review instructions (an 808-row register with 251 verbatim reviewer comments) were re-applied row by row to the 798 baseline L4 cards, producing a 777-card recovery release. Second, an independent audit pipeline (AC-01 through AC-08) was applied to that release. Every judgemental step was performed by two independent reviewer agents with consensus-only application, covering restoration of hard-rule violations, absorption of semantic duplicates, MECE curation of every L3 category (duplication and attribution strength), decomposition of composite cards, and literature-grounded reinforcement of under-populated categories. No EM or Hybrid EM was re-run at any stage. **The final release contains 629 L4 cards: General 487, Agentic 65, Physical 77. All 47 L3 categories hold at least three cards and no Others assignments remain.**

### 1 Recovery principles

Valid edits and user-approved merges from the baseline commit were preserved. Only the deletion, merge, split, reassignment, and definition-revision instructions of the second-round review were applied, traced through `source_row_id` lineage. The recovery stage initially preserved the L3 master; AC-14 subsequently added only the explicitly approved `G_SYS_PERF` category and updated the master hash. A separate fidelity audit of all 251 reviewer comments judged 246 applied, 3 partially applied, and 2 not applied; every affected card carries the verbatim comment and its application result.

### 2 Applied results

### 3 Audit corrections (AC-01 to AC-08)

Following a line-by-line audit of the recovery release and user approval, the corrections below were applied. All judgemental decisions are recorded in the paired reviewer ledgers; only consensus (intersection) verdicts were executed. The second-round human review explicitly

| Action           | Review rows |
|------------------|-------------|
| Reassignment     | 80          |
| Rewrite and keep | 39          |
| Split            | 19          |
| Merge            | 14          |
| Delete           | 14          |

Table 1: 166 approved human-review actions (recovery stage)




Figure 1: Final L4 cards by domain

asked for a new L3 covering deficient performance and capability (two explicit new-L3 requests and seventeen transfer instructions). The request had been held back because it conflicted with the immutability rule for the L3 master; with user approval, AC-14 creates **Performance and Reliability Failure** and reattributes ten cards to it by two-reviewer consensus. The L3 master now holds 47 master categories plus three Others, and its hash has been updated accordingly. All added or revised cards carry verified literature references (XSTest, OR-Bench, AgentHarm, Cao et al. CCS 2019, NIST SP 800-82r3, NIST SP 800-193, Sclar et al. ICLR 2024, Huang et al. ICLR 2024, Moore et al. FAccT 2025). The self-harm/suicide family remains separated into four axes: preparation, execution, encouragement, and crisis-response failure.

## 4 Release composition

The final release ships one L1 CSV, one L1/L2/L3 CSV, and three domain L4 CSVs. Public CSVs carry 28 columns: bilingual hierarchy definitions, facet, act-type, the verbatim human-review comment and its application result, and References; full-column originals including scoring and keyword machinery are preserved in the handover package and archive. The validation folder retains the 808-row instruction register, lineage edges, deletion tombstones, merge/decomposition/c logs, and the audit correction log.

| Correction | Content                                                                            | Cards   |
|------------|------------------------------------------------------------------------------------|---------|
| AC-01      | Restored suicide preparation/execution separation                                  | 777→778 |
| AC-02      | Absorbed 58 consensus duplicates across 8 large L3s                                | 778→720 |
| AC-03      | Annotated verbatim reviewer comments and application results                       | 720     |
| AC-04      | Stripped scoring/keyword columns from public CSVs (full originals archived)        | 720     |
| AC-05      | Populated the two empty L3s with referenced cards                                  | 720→725 |
| AC-06      | Reinforced two under-populated L3s; absorbed 16 governance/manipulation duplicates | 725→711 |
| AC-07      | MECE curation of all L3s: 64 absorbed, 18 reattributed, 2 removed                  | 711→645 |
| AC-08      | Decomposed 18 consensus composite cards                                            | 645→629 |

Table 2: Audit correction summary




Figure 2: Card count across audit corrections